

Local-LLM memory benchmark: 5 tasks $\times$ 5 seeds per method (n=25/method)

A. Strict-exact recall by task and method

Beginning fact survives distractor tail	0.0	1.0	0.0	1.0	1.0	1.0
Stale value overwritten later	0.0	0.0	0.0	1.0	1.0	1.0
Near-duplicate records, multiple entities	0.0	0.0	0.0	1.0	1.0	1.0
Paraphrased query, no lexical match	0.0	0.0	0.0	1.0	1.0	1.0
Many stale values, one authoritative	0.0	0.0	0.0	1.0	1.0	1.0
	Full context	Truncated head	Truncated tail	Keyword RAG	Embedding RAG	Structured slots

B. Cost/accuracy frontier (marker size = strict-exact accuracy)

